

DATABASE MANAGEMENTSYSTEM (MC-3D2)ASSIGNMENT-3

(1)

(i) Write-read conflict

T₁: R(X) R(Y) W(X)T₂: R(X) R(Y) W(X) W(Y)

(ii) Read-write conflict

T₁: R(X) R(Y) W(X)T₂: R(X) R(Y) W(X) W(Y)

(iii) Write-write conflict

T₁: R(X) R(Y) W(X)T₂: R(X) R(Y) W(X) W(Y)(iv) Write-read conflict: T₂ will not get a lock on X until T₁ commits.Read-write conflict: T₁ will not get exclusive lock on X until T₂ commits.Write-write conflict: If T₂ doesn't commit T₁ doesn't get a lock.

(2)

(i) R₁(X) R₂(X) W₁(X) W₂(X)

Not serializable, Not conflict serializable, Not view-serializable, Recoverable, Avoids cascading aborts. NOT strict

(2) $W_1(X) R_2(Y) R_1(Y) R_2(X)$

Serializable, conflict serializable, view serializable,
Does not avoid cascading aborts, NOT strict.

Recoverability can't be decided since commit or
abort sequence is not specified.

(3) $R_1(X) R_1(Y) W_1(X) R_2(Y) W_2(Y) W_1(X) R_2(Y)$



NOT conflict serializable

View and serializability cannot be decided.

NOT avoids cascading aborts

NOT strict

Recoverability cannot be decided since commit
or abort sequence is not specified.

(4) $R_1(X) W_2(X) W_1(X) \text{ Abort}_2 \text{ Commit}_1$

Serializable, conflict serializable & view serializable

Recoverable, avoids cascading aborts

NOT strict

(5) $R_1(X) W_2(X) W_1(X) \text{ Commit}_2 \text{ Commit}_1$

Serializable, NOT Conflict Serializable, view serializable

Recoverable, avoids cascading aborts

NOT strict.

(6) ~~W₁(x) R₂(x) W₁(x) Commit₂ Abort₁~~

Serializable, view serializable, conflict serializable

NOT recoverable, DOES NOT avoid cascading aborts

NOT strict

(7) R₁(x) W₃(x) Commit₃ W₁(y) Commit₁ R₂(y)

W₂(z) Commit₂

Serializable, view serializable, conflict serializable

Recoverable, Avoids cascading aborts, Strict.

(8) R₁(x) W₂(x) W₁(x) R₃(x) Commit₁ Commit₂ Commit₃

Serializable, view serializable, NOT conflict serializable

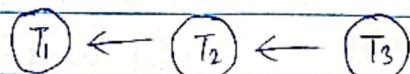
Recoverable, DOES NOT avoid cascading aborts,

NOT strict.

3.

Time	t ₁	t ₂	t ₃	t ₄	t ₅	t ₆	t ₇
T ₁	X(A)						S(C)
T ₂		S(B)	S(C)				S(A)
T ₃		S(C)			X(B)		
L _M	G	G	G	G	B	B	G

(ii) Wait-for graph for the lock requests at time-tick t₇:



(iii) A deadlock does not exist because there is no cycle in the dependency graph.

Q.

ii) Time t_1 t_2 t_3 t_4 t_5 t_6 t_7 t_8

T_1 $X(B)$

$S(A)$

T_2

$X(D)$ $X(C)$

T_3

$S(C)$

T_4

$X(A)$

$X(B)$

ZM

G G G

B

G

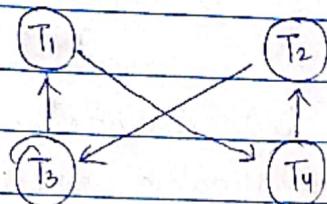
B

$S(D)$

B

B

iii) Wait-for graph for the lock requests.



A deadlock exists because there is a cycle in the dependency graph.

ciii) $X(A)$ at t_2 : G₁

$T_1 < T_2 < T_3 < T_4$

$S(C)$ at t_3 : G₁

$S(A)$ at t_4 : A (T₁ aborts)

$X(D)$ at t_5 : G₁

$X(C)$ at t_6 : A (T₂ aborts)

$X(B)$ at t_7 : G₁ (No locks held on B since T₁ was aborted)

$S(D)$ at t_8 : G₁ (No locks held on D since T₂ was aborted)

civ) $X(A)$ at t_2 : G₁

$S(C)$ at t_3 : G₁

$S(A)$ at t_4 : (T₁ aborts)

$X(D)$ at t_5 : G₁

$X(C)$ at t_6 : B

$X(B)$ at t_7 : G₁ (T_1 aborts)

$S(D)$ at t_8 : G₁ (T_2 aborts)

(5)

B-Tree

(a) Insert: 1, 3, 5, 7, 9, 2, 4, 6, 8, 10

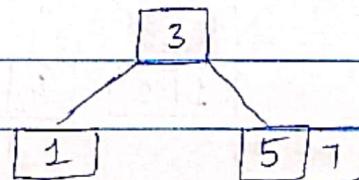
Insert 1:

1

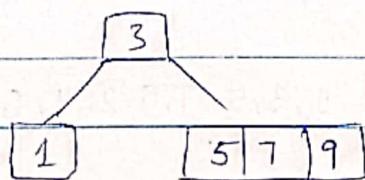
Insert 3, 5:

1 | 3 | 5

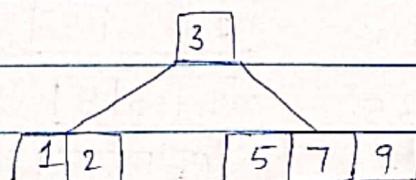
Insert 7:



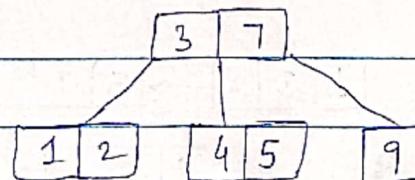
Insert 9:



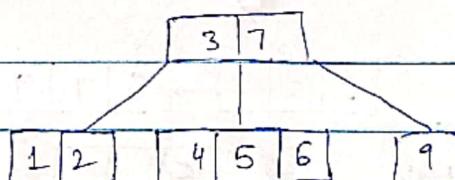
Insert 2:



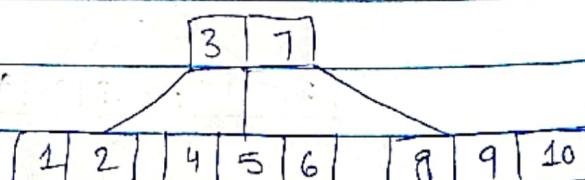
Insert 4:



Insert 6:

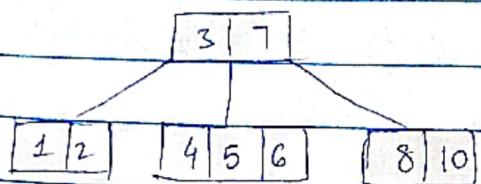


Insert 8 and 10

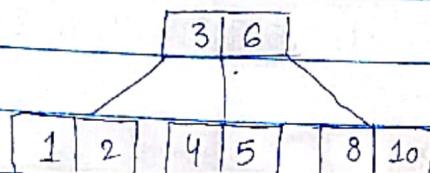


(b) Delete 9, 7, 8

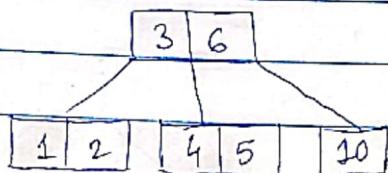
Delete 9:



Delete 7:



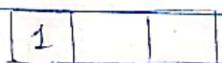
Delete 8:



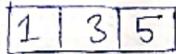
B+ TREES.

(a) Insert 1, 3, 5, 7, 9, 2, 4, 6, 8, 10

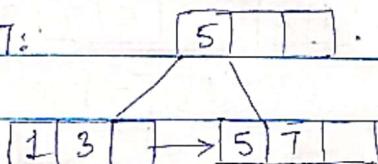
Insert 1:



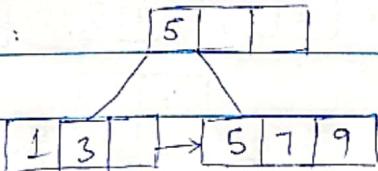
Insert 3, 5:



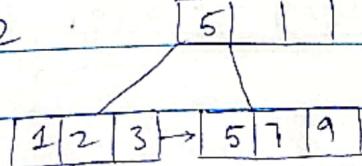
Insert 7:



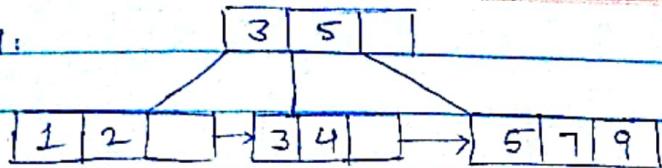
Insert 9:



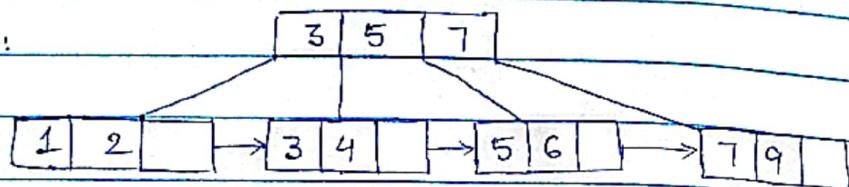
Insert 2:



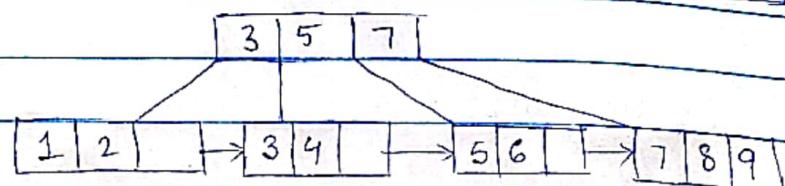
Insert 4:



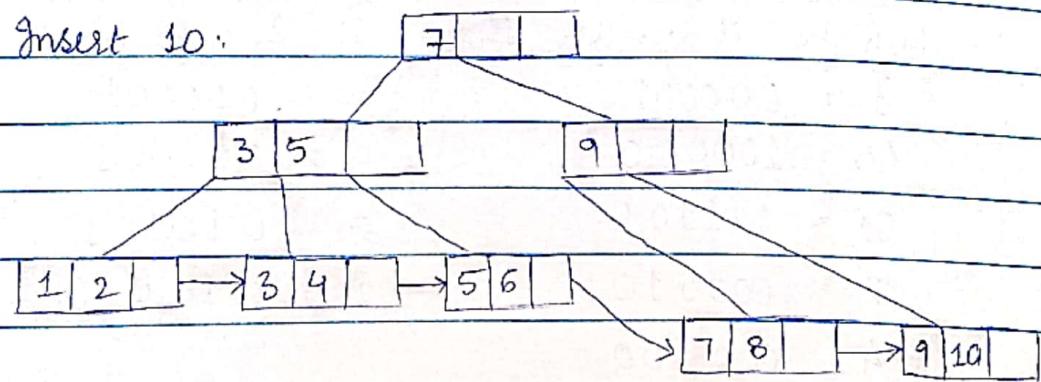
Insert 6:



Insert 8:

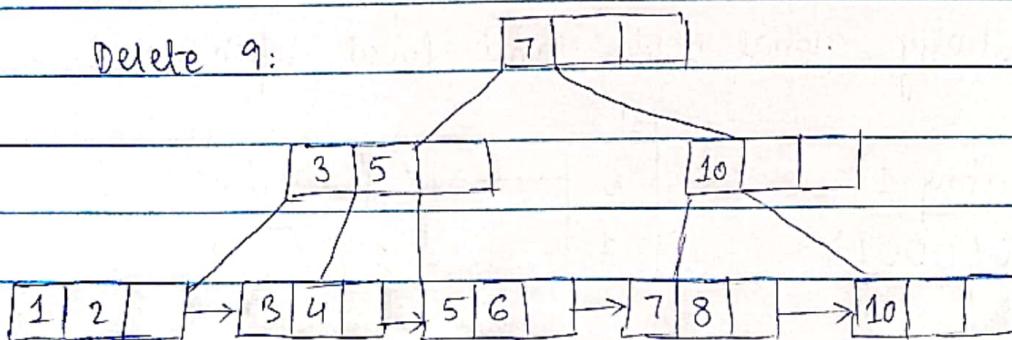


Insert 10:

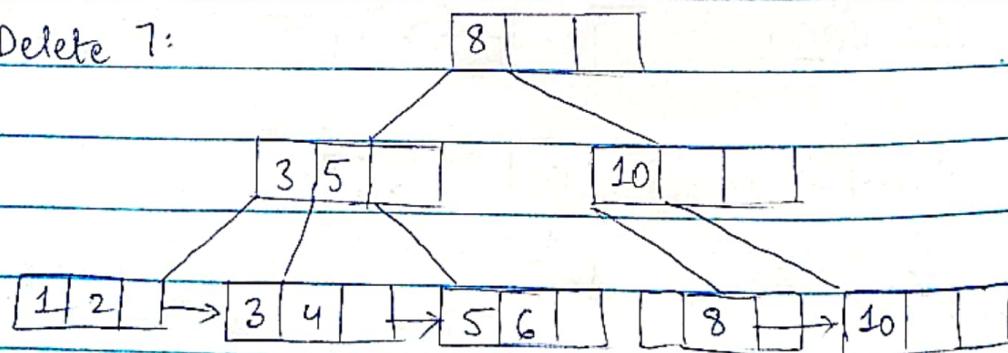


(b) Delete 9, 7, 8

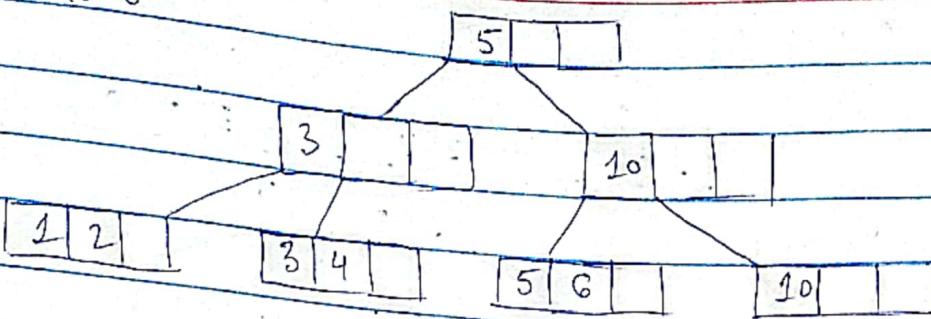
Delete 9:



Delete 7:



Delete 8 :



(6)

Bucket Size = 3.

Hash the elements : 1, 40, 61, 2, 4, 15, 30, 17, 9, 20, 26

$$1 = 000001$$

$$17 = 010001$$

$$40 = 101000$$

$$9 = 001001$$

$$61 = 111101$$

$$20 = 010100$$

$$2 = 000010$$

$$26 = 011010$$

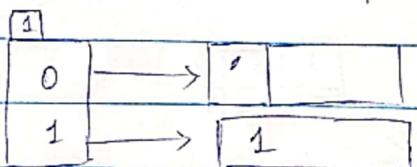
$$4 = 000100$$

$$15 = 001111$$

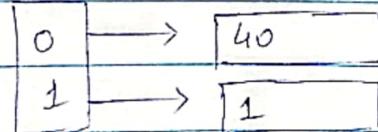
$$30 = 011110$$

Initially global depth and local depth is 1.

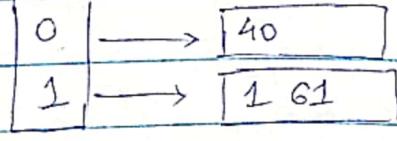
Inserting 1
(000001)



Inserting 40
(101000)



Inserting 61
(111101)



Inserting 2

(0000010)

0	→	40 2
1	→	1 61

Inserting 4 and 15

(0001000) (00111)

0	→	40 2 4
1	→	1 61 15

Inserting 30 : will result in overflow. Directory expansion takes place. Global depth = 2

00	→	40 4
01	→	1 61
10	→	2 30
11	→	15

Inserting 17 :

(010001)

00	→	40 4
01	→	1 61 17
10	→	2 30
11	→	15

Inserting 9 : will result in overflow. Directory expansion takes place. Global Depth = 3

000	→	40
001	→	1 17 9
010	→	2 26
011	→	
100	→	4 20
101	→	61
110	→	30
111	→	15

Inserting 20 and 26

Ans

7.

(a) STATIC HASHING

In static hashing, the resultant data bucket address will always be the same.

For eg. If we generate an address for EMP_ID = 103 using the hash function mod(5) then it will always result in the same bucket address i.e.

Hence in static hashing the number of data buckets in memory remains constant throughout.

(b) Most common hash functions used :

(i) Hashing by Division - The mod method.

The key is mapped into one of the slots by taking the remainder of key divided by table size.

(ii) Hashing by Multiplication.

- The key K is multiplied by a constant real number c in the range $0 < c < 1$ and the fractional part is extracted. ($K * c$)

- This value is multiplied by table size m and floor of the result is taken.

$$h(K) = \text{floor}(m * \text{frac}(K * c))$$

(c) Common collision resolution techniques :

(i) REHASHING

It invokes a secondary hash function which is applied continuously until an empty slot is found.

(ii) CHAINING

It builds a linked list of items whose key hashes to the same value.

(iii) OPEN ADDRESSING

In open addressing we keep probing until an empty slot is found.

It can be done in one of the 3 ways

1. Linear Probing
2. Quadratic Probing
3. Double Hashing.

(d) STATIC HASHING

- 1) A fixed no. of buckets is allocated to a file to store the records
- 2) It uses a fixed hash function to partition the set of all possible search-key values into subsets and then maps each subset to a bucket.

- 3) Performance degrades due to the bucket overflow.

- 4) In static hashing the resultant data bucket address is always the same.

DYNAMIC HASHING

- 1) It allows the number of buckets to vary dynamically.
- 2) It uses a second stage of mapping to determine the bucket associated with some search-key value.

- 3) Performance does not degrade as the file grows.

- 4) In dynamic hashing the data buckets change depending on the records.